

实验三

CIFAR10图像分类

2022年10月2日

1. 读取数据文件

使用torchvision.datasets的CIFAR10数据集，训练集进行旋转和裁剪操作。用DataLoader的方式分批读取

读取数据文件

```
In [2]: train_trans = transforms.Compose([
        transforms.RandomHorizontalFlip(p=0.5), # 随机旋转
        transforms.RandomCrop(32, padding=4), # 放大后随机裁剪到32*32
        transforms.ToTensor(), # 转为tensor
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)) # 归一化
    ])

    test_trans = transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
    ])
    trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=train_trans)
    testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=test_trans)

    Files already downloaded and verified
    Files already downloaded and verified

In [3]: # 用dataloader的方式分批读取
    train_loader = DataLoader(dataset=trainset, batch_size=20, shuffle=True, num_workers=0)
    test_loader = DataLoader(dataset=testset, batch_size=20, shuffle=False, num_workers=0)

    # gpu 加速
    device = torch.device('cpu')
```

实现训练和测试函数

```
In [4]: # 每个epoch的训练结果
    def train_epoch(model, dataloader):
        model.train()
        running_loss = 0
        for i, (img, label) in enumerate(tqdm(dataloader)):
            img, label = img.to(device), label.to(device) # 将数据读入gpu
            out = model(img) # 模型预测结果
            optimizer.zero_grad()
            loss = criterion(out, label) # 计算损失
            loss.backward() # 反向传播更新梯度
            optimizer.step()
            running_loss += loss.item() * img.size(0)
        return running_loss / len(dataloader.dataset)

In [5]: def test(model, dataloader):
        model.eval()
        total = 0
        correct = 0
        for i, (img, label) in enumerate(tqdm(dataloader)):
            img, label = img.to(device), label.to(device)
            out = model(img)
            _, pred = torch.max(out, 1) # 计算正确率
            correct += (pred == label).sum().item()
            total += label.size(0)
        acc = 100.0 * correct / total
        return acc

In [6]: def train(model, n_epochs):
        loss_arr = []
        acc_arr = []
        for epoch in range(n_epochs):
            loss = train_epoch(model, train_loader)
            acc = test(model, train_loader)
            print('Epoch {}, Train Loss {}, Train Acc {}'.format(epoch + 1, loss, acc))
            loss_arr.append(loss)
            acc_arr.append(acc)
        return loss_arr, acc_arr
```


3. CNN模型搭建

搭建了具有3个卷积层和池化层，以及3个全连接层的卷积神经网络

输入的是 $3 * 32 * 32$ 的数据

第一个卷积层使用 $6 * 5 * 5$ 的卷积核，输出 $6 * 28 * 28$

第一个池化层进行 $2 * 2$ 的最大池化，输出 $6 * 14 * 14$

第二个卷积层使用 $16 * 3 * 3$ 的卷积核，输出 $16 * 12 * 12$

第二个池化层进行 $2 * 2$ 的最大池化，输出 $16 * 6 * 6$

第三个卷积层使用 $64 * 2 * 2$ 的卷积核，输出 $64 * 4 * 4$

第三个池化层进行 $2 * 2$ 的最大池化，输出 $64 * 2 * 2$

后面接入三个全连接层，隐含层使用ReLU作为激活函数。

CNN模型搭建

```
In [9]: class CNN(nn.Module):
        def __init__(self):
            super(CNN, self).__init__()
            self.conv1 = nn.Conv2d(3, 6, 5)           # 6 * 28 * 28
            self.pool1 = nn.MaxPool2d(2)            # 6 * 14 * 14
            self.conv2 = nn.Conv2d(6, 16, 3)         # 16 * 12 * 12
            self.pool2 = nn.MaxPool2d(2)            # 16 * 6 * 6
            self.conv3 = nn.Conv2d(16, 64, 3)        # 64 * 4 * 4
            self.pool3 = nn.MaxPool2d(2)            # 64 * 2 * 2
            self.fc1 = nn.Linear(64 * 2 * 2, 1024)
            self.fc2 = nn.Linear(1024, 256)
            self.fc3 = nn.Linear(256, 10)

        def forward(self, x):
            x = self.pool1(F.relu(self.conv1(x)))
            x = self.pool2((self.conv2(x)))
            x = self.pool3(F.relu(self.conv3(x)))
            x = x.view(-1, 64 * 2 * 2)
            x = F.relu(self.fc1(x))
            x = F.relu(self.fc2(x))
            x = self.fc3(x)
            return x
```

使用CrossEntropyLoss作为损失函数，使用学习率为0.01，衰退权重为

0.0001的SGD优化器进行训练。50次迭代后的预测正确率为74.03%。

```
In [10]: cnn = CNN().to(device)
criterion = nn.CrossEntropyLoss().to(device)
optimizer = torch.optim.SGD(params=cnn.parameters(), lr=0.01, weight_decay=0.0001)
cnn_loss, cnn_acc = train(cnn, n_epochs)
print('Test Acc: {}'.format(test(cnn, test_loader)))
```

```
100%|████████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:18<00
:00, 137.19it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:08<00
:00, 306.63it/s]
```

Epoch 49, Train Loss 0.6774460505485534, Train Acc 76.284

```
100%|████████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:18<00
:00, 136.23it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:08<00
:00, 307.28it/s]
```

Epoch 50, Train Loss 0.6712483007788658, Train Acc 77.026

```
100%|████████████████████████████████████████████████████████████████████████████████| 500/500 [00:01<00
:00, 374.38it/s]
```

Test Acc: 74.03


3. 绘制变化曲线

使用前两个实验的代码

```
import matplotlib.pyplot as plt
def plot_loss_and_acc(loss_and_acc_dict):
    # visualize loss curve
    plt.figure(figsize=(12, 5))
    min_loss, max_loss = 100.0, 0.0
    for key, (loss_list, acc_list) in loss_and_acc_dict.items():
        min_loss = min(loss_list) if min(loss_list) < min_loss else min_loss
        max_loss = max(loss_list) if max(loss_list) > max_loss else max_loss
        num_epoch = len(loss_list)
        plt.plot(range(1, 1 + num_epoch), loss_list, '-s', label=key)
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()
    plt.xticks(range(0, num_epoch + 1, 2))
    plt.axis([0, num_epoch + 1, min_loss - 0.1, max_loss + 0.1])
    plt.show()
    # visualize acc curve
    plt.figure(figsize=(12, 5))
    min_acc, max_acc = 100.0, 0.0
    for key, (loss_list, acc_list) in loss_and_acc_dict.items():
        min_acc = min(acc_list) if min(acc_list) < min_acc else min_acc
        max_acc = max(acc_list) if max(acc_list) > max_acc else max_acc
        num_epoch = len(acc_list)
        plt.plot(range(1, 1 + num_epoch), acc_list, '-s', label=key)
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.xticks(range(0, num_epoch + 1, 2))
    plt.axis([0, num_epoch + 1, min_acc, max_acc + 1])
    plt.show()
```

从结果中可以看出，CNN的loss更小，而且正确率也比MLP的要高20%

```
plot_loss_and_acc({
    "mlp": [mlp_loss, mlp_acc],
    "cnn": [cnn_loss, cnn_acc]
})
```



4. 调整优化器

使用学习率为0.001的Adam优化器训练

```
cnn = CNN().to(device)
optimizer = torch.optim.Adam(params=cnn.parameters(), lr=0.001)
cnn_adam_loss, cnn_adam_acc = train(cnn, n_epochs)

:00, 110.89it/s]
100% |██████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:08<00
:00, 298.77it/s]

Epoch 48, Train Loss 0.7755873042285443, Train Acc 73.206

100% |██████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:22<00
:00, 110.44it/s]
100% |██████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:08<00
:00, 298.59it/s]


Epoch 49, Train Loss 0.7645300450325012, Train Acc 73.874

100% |██████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:22<00
:00, 111.77it/s]
100% |██████████████████████████████████████████████████████████████████████████████| 2500/2500 [00:08<00
:00, 299.41it/s]

Epoch 50, Train Loss 0.7688073013603687, Train Acc 73.496
```

与使用SGD优化器对比为：

```
plot_loss_and_acc({
    "sgd": [cnn_loss, cnn_acc],
    "adam": [cnn_adam_loss, cnn_adam_acc]
})
```



可以看出一开始Adam优化器的损失下降的较快，但是随着epoch的增加，adam的损失则变化得较小，而SGD的正确率也逐渐超过Adam。

5. 调整网络层数

增加一个全连接层，并在全连接层之间添加了一个dropout层。

```
class CNNNew(nn.Module):
    def __init__(self):
        super(CNNNew, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)          # 6 * 28 * 28
        self.pool1 = nn.MaxPool2d(2)             # 6 * 14 * 14
        self.conv2 = nn.Conv2d(6, 16, 3)          # 16 * 12 * 12
        self.pool2 = nn.MaxPool2d(2)             # 16 * 6 * 6
        self.conv3 = nn.Conv2d(16, 64, 3)         # 64 * 4 * 4
        self.pool3 = nn.MaxPool2d(2)             # 64 * 2 * 2
        self.fc1 = nn.Linear(64 * 2 * 2, 2048)
        self.fc2 = nn.Linear(2048, 512)
        self.fc3 = nn.Linear(512, 256)
        self.fc4 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool1(F.relu(self.conv1(x)))
        x = self.pool2(self.conv2(x))
        x = self.pool3(F.relu(self.conv3(x)))
        x = x.view(-1, 64 * 2 * 2)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.dropout(x, 0.3)                    # 增加dropout层
        x = F.relu(self.fc3(x))
        x = self.fc4(x)
        return x
```

使用CrossEntropyLoss作为损失函数，使用学习率为0.01，衰退权重为0.0001的SGD优化器进行训练。50次迭代后的预测正确率为73.13%。

```
cnn = CNNNew().to(device)
criterion = nn.CrossEntropyLoss().to(device)
optimizer = torch.optim.SGD(params=cnn.parameters(), lr=0.01, weight_decay=0.0001)
cnn_new_loss, cnn_new_acc = train(cnn, n_epochs)
print('Test Acc: {}'.format(test(cnn, test_loader)))
```

100%|██| 2500/2500 [00:27<00:00, 91.29it/s]

100%|██| 2500/2500 [00:10<00:00, 244.78it/s]

Epoch 49, Train Loss 0.6798847847402095, Train Acc 76.496

100%|██| 2500/2500 [00:27<00:00, 91.15it/s]

100%|██| 2500/2500 [00:10<00:00, 243.20it/s]


Epoch 50, Train Loss 0.6801970268666744, Train Acc 75.998

100%|██| 500/500 [00:01<00:00, 283.65it/s]

Test Acc: 73.13

与原来CNN的对比中，可以看到两者曲线基本一致，并没有看出增加层数带来的效果。

```
plot_loss_and_acc({
    "cnn": [cnn_loss, cnn_acc],
    "new_cnn": [cnn_new_loss, cnn_new_acc]
})
```



6. 修改隐含层激活函数

将隐含层激活函数替换为Tanh

```
class CNNTanh(nn.Module):
    def __init__(self):
        super(CNNTanh, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)      # 6 * 28 * 28
        self.pool1 = nn.MaxPool2d(2)         # 6 * 14 * 14
        self.conv2 = nn.Conv2d(6, 16, 3)     # 16 * 12 * 12
        self.pool2 = nn.MaxPool2d(2)         # 16 * 6 * 6
        self.conv3 = nn.Conv2d(16, 64, 3)    # 64 * 4 * 4
        self.pool3 = nn.MaxPool2d(2)         # 64 * 2 * 2
        self.fc1 = nn.Linear(64 * 2 * 2, 1024)
        self.fc2 = nn.Linear(1024, 256)
        self.fc3 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool1(F.relu(self.conv1(x)))
        x = self.pool2((self.conv2(x)))
        x = self.pool3(F.relu(self.conv3(x)))
        x = x.view(-1, 64 * 2 * 2)
        x = F.tanh(self.fc1(x))
        x = F.tanh(self.fc2(x))
        x = self.fc3(x)
        return x
```



```

cnn = CNNTanh().to(device)
criterion = nn.CrossEntropyLoss().to(device)
optimizer = torch.optim.SGD(params=cnn.parameters(), lr=0.01, weight_decay=0.0001)
cnn_tanh_loss, cnn_tanh_acc = train(cnn, n_epochs)
print('Test Acc: {}'.format(test(cnn, test_loader)))

```

100%|██| 2500/2500 [00:21<00:00, 118.18it/s]

100%|██| 2500/2500 [00:08<00:00, 280.71it/s]

Epoch 49, Train Loss 0.7560197077989578, Train Acc 74.762

100%|██| 2500/2500 [00:21<00:00, 117.66it/s]

100%|██| 2500/2500 [00:08<00:00, 280.30it/s]

Epoch 50, Train Loss 0.7479605409920216, Train Acc 74.364

100%|██| 500/500 [00:01<00:00, 337.27it/s]

Test Acc: 73.51

对比中，可以看出，ReLU作为激活函数的效果要更好一点。

```

plot_loss_and_acc({
    "relu": [cnn_loss, cnn_acc],
    "tanh": [cnn_tanh_loss, cnn_tanh_acc]
})

```

